Retrieval Optimization with Semantic Cache

Ziming Wang, Tim Han, Dingfan Zheng, Lei Li
{peterwa2,timhan,dingfanz,leili}@cs.cmu.edu

Carnegie Mellon University

Caching in LLM RAG Applications

- Previous work caches LLM outputs, not fine-grained enough
- Semantically similar question retrieve similar documents, cache the query and <u>retrieved document</u> instead!

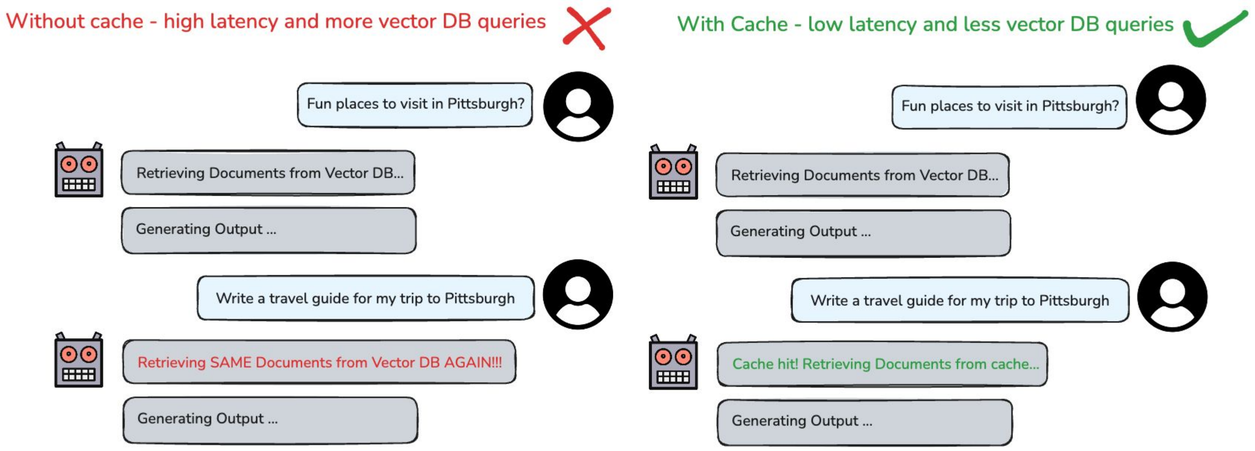

Semantic Cache System Design

- Reuse topk results for similar queries based on distance threshold.
- Tunable cache size, threshold, and eviction policies for optimization.

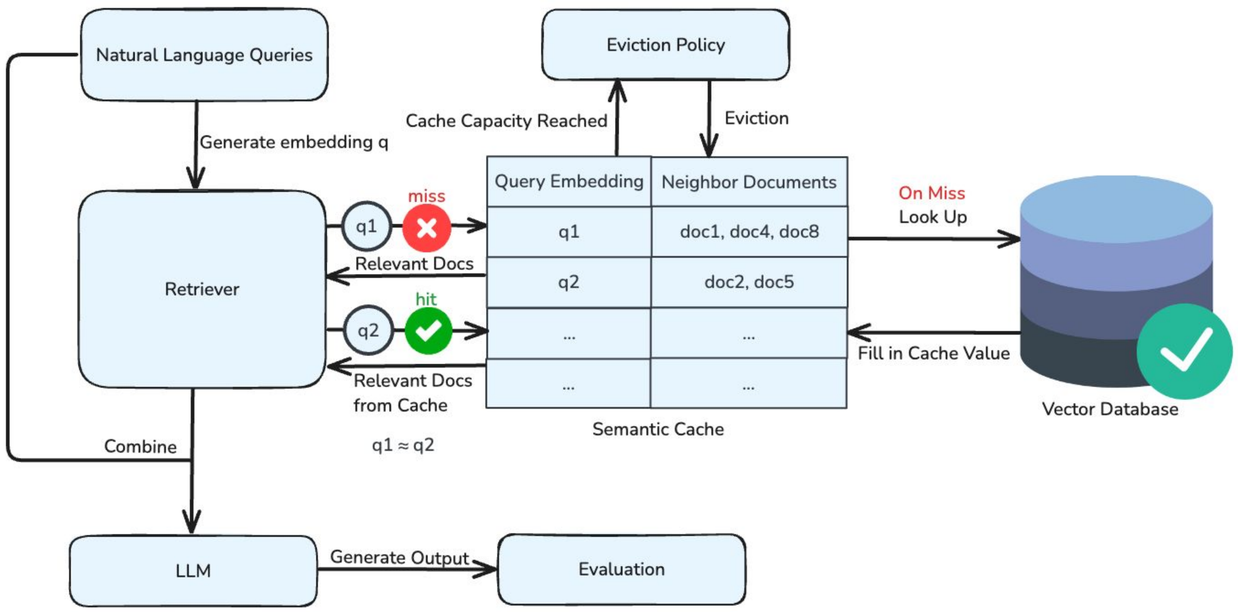

Results

- Our approach cuts retrieval latency by > 30%, while maintaining high retrieval quality with only a 2–4% drop in accuracy.

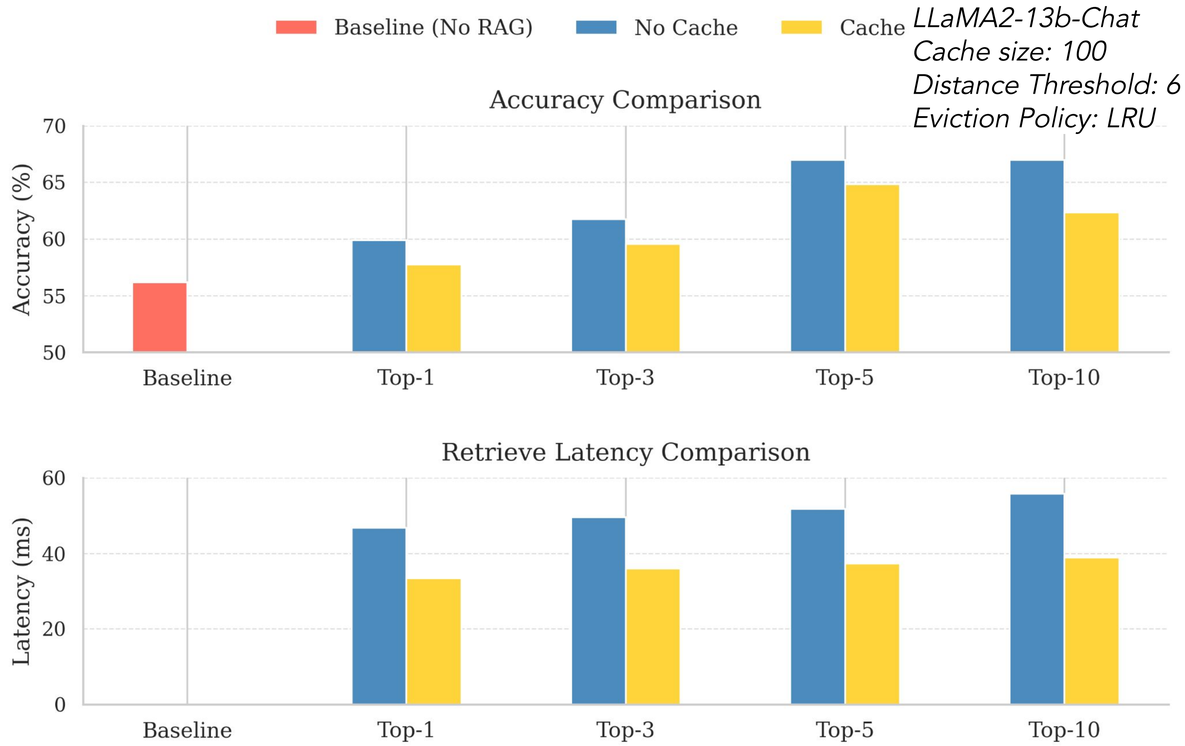

- Impact of Cache Size and Threshold on Performance Metrics

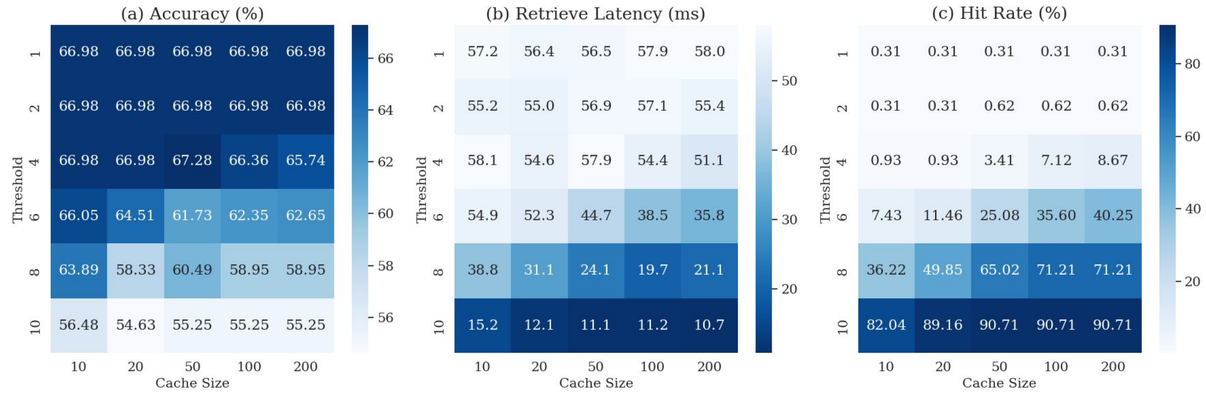

Key Takeaway / Advantages

- Reduces retrieval latency and the number of API calls
- Tuning cache configs based on workloads and resources